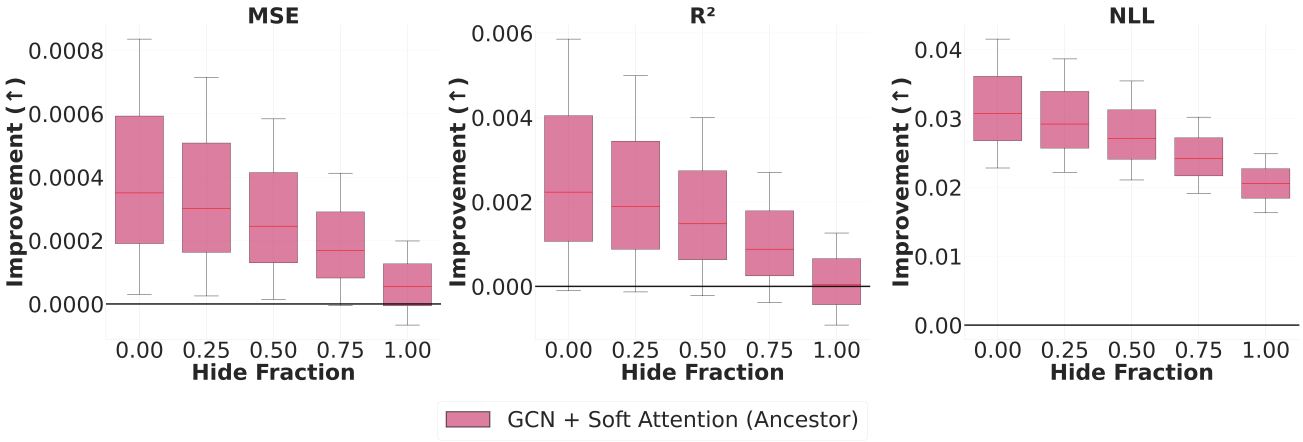

Figure 1. Performance on complex synthetic data for different fractions of hidden entries in the ancestor matrix for the GCN + Soft Attention (Ancestor) graph-conditioning architecture trained with graphical information. Results are measured as the mean of the absolute differenced to a model that uses the same GCN component, but is trained and evaluated **without** graph information. The difference to Figure 4 in the original paper is that the baseline now has this GCN component. The error-bars represent 95-percent bootstrap confidence intervals for the mean. Providing full graph information significantly improves performance, even when fixing the architecture.